

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df=pd.read_csv("Churn_Modelling.csv")
df
```

```
Out[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1
...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1

10000 rows × 14 columns



```
In [4]: df.isnull().sum()
```

```
Out[4]: RowNumber      0
CustomerId    0
Surname       0
CreditScore   0
Geography     0
Gender        0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
```

```
In [5]: df.info
```

```
Out[5]: <bound method DataFrame.info of
0      1  15634602  Hargrave      619  France  Female  42
1      2  15647311    Hill      608   Spain  Female  41
2      3  15619304   Onio      502   France  Female  42
3      4  15701354   Boni      699   France  Female  39
4      5  15737888  Mitchell     850   Spain  Female  43
...    ...    ...    ...    ...    ...    ...
9995   9996  15606229  Obijiaku     771   France   Male  39
9996   9997  15569892  Johnstone   516   France   Male  35
9997   9998  15584532    Liu      709   France  Female  36
9998   9999  15682355  Sabbatini   772  Germany   Male  42
9999  10000  15628319   Walker     792   France  Female  28
```

```

      Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0          2    0.00           1          1          1
1          1  83807.86           1          0          1
2          8 159660.80           3          1          0
3          1    0.00           2          0          0
4          2 125510.82           1          1          1
...    ...    ...    ...    ...    ...
9995       5    0.00           2          1          0
9996      10  57369.61           1          1          1
9997       7    0.00           1          0          1
9998       3   75075.31           2          1          0
9999       4 130142.79           1          1          0
```

```

      EstimatedSalary  Exited
0          101348.88        1
1          112542.58        0
2          113931.57        1
3           93826.63        0
4           79084.10        0
...    ...    ...
9995       96270.64        0
9996      101699.77        0
9997       42085.58        1
9998       92888.52        1
9999       38190.78        0
```

```
[10000 rows x 14 columns]>
```

```
In [6]: df.dtypes
```

```
Out[6]: RowNumber      int64
CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age            int64
Tenure          int64
Balance         float64
NumOfProducts  int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
Exited          int64
dtype: object
```

```
In [7]: df.columns
```

```
Out[7]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [8]: df.head()
```

```
Out[8]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1			
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0			
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1			
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0			
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1			


```
In [9]: df=df.drop(['RowNumber','CustomerId','Surname'],axis=1)
```

```
In [10]: def visualization(x,y,xlabel):
```

```
plt.figure(figsize=(10,5))
plt.hist([x,y],color=['red','green'],label=['exit','not_exit'])
plt.xlabel(xlabel,fontsize=20)
plt.ylabel("No. of cusstomers",fontsize=20)
plt.legend()
```


```
In [11]: df_churn_exited=df[df['Exited']==1]['Tenure']
df_churn_not_exited=df[df['Exited']==0]['Tenure']
```

```
In [12]: visualization(df_churn_exited,df_churn_not_exited,"Tenure")
```



```
In [13]: df_churn_exited=df[df['Exited']==1]['Age']
df_churn_not_exited=df[df['Exited']==0]['Age']
```

```
In [14]: visualization(df_churn_exited,df_churn_not_exited,"Age")
```



```
In [15]: df.dtypes
```

```
Out[15]: CreditScore      int64
Geography      object
Gender          object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object
```

```
In [16]: x = df[['CreditScore','Gender','Age','Tenure','Balance', 'NumOfProducts','HasCrCard']]
states = pd.get_dummies(df[['Geography']],drop_first= True)
gender = pd.get_dummies(df[['Gender']],drop_first= True)
```

```
In [18]: df = pd.concat([df,gender,states], axis=1)
```

```
In [19]: df.head()
```

Out[19]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.88	0
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	0
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [20]: x = df[['CreditScore','Age','Tenure','Balance', 'NumOfProducts','HasCrCard','IsActiveMember',
               'EstimatedSalary','Germany','Spain','Male']]
y = df[['Exited']]
```

```
In [21]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,random_state=0, test_size=0.25)
```

```
In [22]: x.shape
```

```
Out[22]: (10000, 11)
```

```
In [24]: x_test.shape
```

```
Out[24]: (2500, 11)
```

```
In [25]: x_train
```

Out[25]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Germany	Spain	Male
2967	579	39	5	117833.30	3	0	0	5831.00	True	False	False
700	750	32	5	0.00	2	1	0	95611.47	False	False	False
3481	729	34	9	53299.96	2	1	1	42855.97	False	True	False
1621	689	38	5	75075.14	1	1	1	8651.92	False	True	True
800	605	52	7	0.00	2	1	1	173952.50	False	False	True
...
9225	594	32	4	120074.97	2	1	1	162961.79	True	False	False
4859	794	22	4	114440.24	1	1	1	107753.07	False	True	False
3264	738	35	5	161274.05	2	1	0	181429.87	False	False	True
9845	590	38	9	0.00	2	1	1	148750.16	False	True	False
2732	623	48	1	108076.33	1	1	0	118855.26	True	False	False

7500 rows × 11 columns

```
In [26]: x_train.shape
```

```
Out[26]: (7500, 11)
```

```
In [28]: from sklearn.neural_network import MLPClassifier
ann = MLPClassifier(hidden_layer_sizes=(100,100,100),
                    random_state =0,
```

```
max_iter=100, activation='relu')
ann.fit(x_train,y_train)
```

```
C:\Users\Welcome\AppData\Local\anaconda3\Lib\site-packages\sklearn\network\_multilayer_perceptron.py:690:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged
yet.
  warnings.warn(
```

```
Out[28]: ▼ MLPClassifier
MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100, random_state=0)
```

```
In [29]: y_pred = ann.predict(x_test)
y_pred
```

```
Out[29]: array([0, 0, 0, ..., 0, 1, 0], dtype=int64)
```


```
In [30]: import sklearn
from sklearn.metrics import ConfusionMatrixDisplay, classification_report
from sklearn.metrics import accuracy_score
```

```
In [31]: y_test.value_counts()
```

```
Out[31]: Exited
0      1991
1       509
Name: count, dtype: int64
```

```
In [32]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[32]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1dbdc3cade0>
```



```
In [33]: accuracy_score(y_test,y_pred)
```

```
Out[33]: 0.6888
```

```
In [34]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.79	0.80	1991
1	0.26	0.28	0.27	509
accuracy			0.69	2500
macro avg	0.53	0.54	0.53	2500
weighted avg	0.70	0.69	0.69	2500

```
In [ ]:
```